

Task4_ALI


December 6, 2022

```
[2]: import pandas as pd
import numpy as np
from sklearn import metrics
from sklearn.cluster import KMeans
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
import scipy.stats as stats
import sklearn
from sklearn.cluster import DBSCAN
from sklearn.neighbors import NearestNeighbors
from collections import Counter
from sklearn.preprocessing import StandardScaler
```

```
[3]: df2 = pd.read_csv("Shuttle22.csv", header= None)
df = pd.read_csv("complex9_gn8.txt", header = None, sep=',')
```

```
[18]: def purity_score(y_true, y_pred, outliers):
    # compute contingency matrix (also called confusion matrix)
    contingency_matrix = metrics.cluster.contingency_matrix(y_true, y_pred)
    # return purity
    return np.sum(np.amax(contingency_matrix, axis=0)) / np.
    ↪sum(contingency_matrix)
```

```
[87]: X=df[[0,1]]
X=X.values
X
sns.scatterplot(X[:,0], X[:, 1])
data = df[[0,1]]
plt.show()
data,X
```



```
[87]: (
      0      1
0    660.976 304.2250
1    636.213 306.1740
2    662.753 307.5650
3    657.487 307.7400
4    635.273 308.1570
...
3268 728.899 535.6270
3269 504.528 -46.2297
3270 373.256 409.0260
3271 850.838 242.7110
3272 641.676 347.5440

[3273 rows x 2 columns],
array([[660.976, 304.225],
       [636.213, 306.174],
       [662.753, 307.565],
       ...,
       [373.256, 409.026],
       [850.838, 242.711],
       [641.676, 347.544]]))
```

```
[102]: def calculate_cost(X, centroids, cluster):
        sum = 0
        for i, val in enumerate(X):
```

```


        sum += np.sqrt((centroids[int(cluster[i]), 0]-val[0])**2 +
↪(centroids[int(cluster[i]), 1]-val[1])**2)
    return sum
def kmeans(X, k):
    diff = 1
    cluster = np.zeros(X.shape[0])
    centroids = data.sample(n=k).values
    while diff:
        # for each observation
        for i, row in enumerate(X):
            mn_dist = float('inf')
            # dist of the point from all centroids
            for idx, centroid in enumerate(centroids):
                d = np.sqrt((centroid[0]-row[0])**2 + (centroid[1]-row[1])**2)
                # store closest centroid
                if mn_dist > d:
                    mn_dist = d
                    cluster[i] = idx
            new_centroids = pd.DataFrame(X).groupby(by=cluster).mean().values
        # if centroids are same then leave
        if np.count_nonzero(centroids-new_centroids) == 0:
            diff = 0
        else:
            centroids = new_centroids
    return centroids, cluster
cost_list = []
for k in range(1, 10):
    centroids, cluster = kmeans(X, k)
    # WCSS (Within cluster sum of square)
    cost = calculate_cost(X, centroids, cluster)
    cost_list.append(cost)

```

```

[104]: sns.lineplot(x=range(1,10), y=cost_list, marker='o')
plt.xlabel('k')
plt.ylabel('WCSS')
plt.show()


```



```
[143]: palette=sns.color_palette("viridis", as_cmap=True)

k = 9
centroids, cluster = kmeans(X, k)

sns.scatterplot(X[:,0], X[:, 1], hue=cluster,palette=palette)
sns.scatterplot(centroids[:,0], centroids[:, 1], s=100, color='r')
plt.show()
```




```
[154]: true = df[2]
       purity_score(true,cluster)
```

```
[154]: 0.7100519401161014
```

```
[155]: k = 13
       centroids, cluster = kmeans(X, k)

       sns.scatterplot(X[:,0], X[:, 1], hue=cluster,palette=palette)
       sns.scatterplot(centroids[:,0], centroids[:, 1], s=100, color='r')
       plt.show()
       purity_score(true,cluster)
```



[155]: 0.6804155209288115

```
[4]: z = df2[[0,1,2,3,4,5,6,7,8]]
ZSHUT=z.apply(stats.zscore)
ZSHUT[9] = df2[9]

zdata = ZSHUT[[0,1,2,3,4,5,6,7,8]]
zdata= zdata.values
zdata
#z
```

```
[4]: array([[ 1.43954104e-01,  2.69627492e-01, -9.37819783e-01, ...,
            -7.69740511e-01, -1.34679698e-01,  3.14970333e-01],
            [ 5.52518390e-01,  2.49473239e-04,  7.47063843e-01, ...,
            -8.33105101e-02,  1.91967997e+00,  1.64237992e+00],
            [ 3.89092675e-01,  2.49473239e-04, -3.76191907e-01, ...,
            -6.17200511e-01, -9.75099563e-01, -4.65858839e-01],
            ...,
            [ 5.52518390e-01,  2.49473239e-04, -9.37819783e-01, ...,
            -1.15109051e+00,  6.59050174e-01,  1.09579950e+00],
            [-9.18313039e-01,  2.49473239e-04,  1.98264517e+00, ...,
             2.20478949e+00,  1.59285002e+00,  2.36887415e-01],
            [ 6.34231247e-01,  2.59045226e-02,  1.42101729e+00, ...,
             3.74309491e-01, -2.28059683e-01, -3.87775921e-01]])
```

```
[5]: data =ZSHUT[[0,1,2,3,4,5,6,7,8]]
data
```

```
[5]:
```

	0	1	2	3	4	5	6 \
0	0.143954	0.269627	-0.937820	-0.00711	-0.302395	-0.007391	-0.769741
1	0.552518	0.000249	0.747064	-0.00711	-1.595103	0.112097	-0.083311
2	0.389093	0.000249	-0.376192	-0.00711	0.805641	-0.030369	-0.617201
3	-0.918313	0.000249	-1.050145	-0.00711	-0.302395	0.075331	0.221769
4	-0.918313	0.000249	-0.713169	-0.00711	-0.025386	-0.126878	0.450579
...
57995	2.595340	0.000249	-0.151541	-0.00711	-3.257157	-0.140665	-2.523951
57996	0.552518	0.000249	-0.488517	-0.00711	-2.518466	0.107501	-0.846011
57997	0.552518	0.000249	-0.937820	-0.00711	-1.041086	-0.108496	-1.151091
57998	-0.918313	0.000249	1.982645	-0.00711	-0.764077	-0.080922	2.204789
57999	0.634231	0.025905	1.421017	-0.00711	0.805641	-0.002795	0.374309
		7	8				
0	-0.13468	0.314970					
1	1.91968	1.642380					
2	-0.97510	-0.465859					
3	-0.13468	-0.231610					
4	-0.22806	-0.465859					
...					
57995	3.22700	3.984867					
57996	2.38658	2.423209					
57997	0.65905	1.095800					
57998	1.59285	0.236887					
57999	-0.22806	-0.387776					

[58000 rows x 9 columns]


```
[353]: k = 3

centroids, cluster = kmeans(zdata, k)

true2 = ZSHUT[9]
purity_score(true2,cluster)
#cluster
```


```
[353]: 0.868948275862069
```

```
[340]: sns.scatterplot(zdata[:,0], zdata[:, 1])
plt.show()
```




```
[362]: c=data
c[9] = cluster.tolist()
c.loc[c[9]==1]
box0 =c.loc[c[9]==0]
box1 =c.loc[c[9]==1]
box2 =c.loc[c[9]==2]
box0=box0[[0,1,2,3,4,5,6,7,8]]
box1=box1[[0,1,2,3,4,5,6,7,8]]
box2=box2[[0,1,2,3,4,5,6,7,8]]
```


```
[369]: plt.figure(figsize=(20,15))
plt.boxplot(box0.values,showfliers=False,
            flierprops={'marker': 'o', 'markersize': 1, 'markerfacecolor': 'fuchsia'})
plt.show()
```


```
[365]: plt.figure(figsize=(20,15))
plt.boxplot(box1.values,showfliers=False,
            flierprops={'marker': 'o', 'markersize': 1, 'markerfacecolor': 'fuchsia'})
plt.show()
```



```
[366]: plt.figure(figsize=(20,15))
plt.boxplot(box2.values,showfliers=False,
            flierprops={'marker': 'o', 'markersize': 1, 'markerfacecolor': 'fuchsia'})
plt.show()
```



```
[373]: plt.figure(figsize=(20,15))
plt.boxplot(data.values,showfliers=False,
            flierprops={'marker': 'o', 'markersize': 1, 'markerfacecolor': 'fuchsia'})
plt.show()
```



```
[64]: def dbscan(dataset, epsilon, minsamples):
    dbscan_data = dataset[[0,1,2,3,4,5,6,7,8,9]]
    model1 = DBSCAN(eps=epsilon, min_samples=minsamples, metric='euclidean').\
    fit(dbscan_data)

    outliers_df = dataset[model1.labels_ == -1]
    num_clusters = len(set(model1.labels_))

    clusters_df = [dataset[model1.labels_ == n] for n in range(num_clusters)]

    colors = model1.labels_
    color_clusters = colors[colors != -1]
    color_outliers = 'white'
    clusters1 = Counter(model1.labels_)
    return clusters_df

def dbscanplot(dataset, epsilon, minsamples):
    dbscan_data = dataset[[0,1,2,3,4,5,6,7,8,9]]

    #return dbscan_data
    model1 = DBSCAN(eps=epsilon, min_samples=minsamples, metric='euclidean').\

```

```

fit(dbscan_data)
#return model1
outliers_df = dataset[model1.labels_ == -1]
clusters_df = dataset[model1.labels_ != -1]
return model1.labels_
'''colors = model1.labels_
color_clusters = colors[colors != -1]
color_outliers = 'white'
clusters1 = Counter(model1.labels_)
print(clusters_df)
print(clusters1)
#print(dataset[model1.labels_ == -1].head())
print('number of clusters: {}'.format(len(clusters1)-1))
#dbscan_plot(clusters_df,color_clusters)'''
def dbscan_plot(cluster,cluster_colors):
fig = plt.figure()
ax = fig.add_axes([.2,.2,2,2])
ax.scatter(ZSHUT[[0,1,2,3,4,5,6,7,8]],
c = cluster_colors, edgecolors = 'black', s = 70)
ax.set_xlabel('Latitude', fontsize=10)
ax.set_ylabel('Longitude', fontsize=10)
plt.title('title',fontsize=12)
plt.grid(which='major',color='#cccccc', alpha=0.45)
plt.show()

```

```

[8]: dbscandata= ZSHUT[[0,1,2,3,4,5,6,7,8]]
dbscanplot(dbscandata,.26,100)

```

```

[8]: array([-1,  0,  1, ...,  3,  5,  1], dtype=int64)

```

```

[9]: clusterlist =dbscan(dbscandata,.26,100)

```

```

[11]: q = dbscanplot(dbscandata,.26,100)

```

```

[36]: n=np.delete(q, np.where(q == -1))
ZSHUT[9].max

```

```

[36]: <bound method NDFrame._add_numeric_operations.<locals>.max of 0      2
1          4
2          1
3          1
4          1
..
57995      5
57996      4
57997      4
57998      1

```

```
57999    4
Name: 9, Length: 58000, dtype: int64>
```

```
[40]: clusterlist[1]
```

```
[40]:
```

	0	1	2	3	4	5	6	\
2	0.389093	0.000249	-0.376192	-0.007110	0.805641	-0.030369	-0.617201	
3	-0.918313	0.000249	-1.050145	-0.007110	-0.302395	0.075331	0.221769	
4	-0.918313	0.000249	-0.713169	-0.007110	-0.025386	-0.126878	0.450579	
7	0.552518	-0.012578	1.084041	-0.089254	0.897977	-0.025773	0.221769	
10	-0.918313	0.000249	-0.825494	-0.061873	-1.041086	-0.007391	0.374309	
...	
57990	-0.836600	0.025905	-0.713169	-0.007110	0.159287	0.075331	0.374309	
57992	-0.754887	-0.025406	-0.600843	-0.116636	0.159287	-0.007391	0.298039	
57993	-0.428036	0.000249	-0.488517	0.020271	0.343959	-0.048752	-0.007041	
57994	0.062241	0.000249	0.185436	-0.007110	0.528632	-0.062539	0.069229	
57999	0.634231	0.025905	1.421017	-0.007110	0.805641	-0.002795	0.374309	
		7	8					
2		-0.97510	-0.465859					
3		-0.13468	-0.231610					
4		-0.22806	-0.465859					
7		-0.46151	-0.465859					
10		0.65905	0.393053					
...						
57990		-0.46151	-0.543942					
57992		-0.46151	-0.543942					
57993		-0.55489	-0.465859					
57994		-0.46151	-0.465859					
57999		-0.22806	-0.387776					

[42135 rows x 9 columns]

```
[51]: ze = ZSHUT
ze[10]=q
ze2 = ze[ze[10] != -1]

truer = ze2[9]
mine = ze2[10]
purity_score(truer,mine)
```

```
[51]: 0.9325662054264966
```

```
[4]: def dbscanplot1(dataset, epsilon, minsamples):
      dbscan_data = dataset[[0,1,2,3,4,5,6,7,8]]

      #return dbscan_data
```

```

model1 = DBSCAN(eps=epsilon, min_samples=minsamples, metric='euclidean').\
fit(dbscan_data)
#return model1
outliers_df = dataset[model1.labels_ == -1]
clusters_df = dataset[model1.labels_ != -1]
colors = model1.labels_
color_clusters = colors[colors != -1]
color_outliers = 'white'
clusters1 = Counter(model1.labels_)
#return model1.labels_
#print(clusters_df)
print(clusters1)
#print(dataset[model1.labels_ == -1].head())
print('number of clusters: {}'.format(len(clusters1)-1))
#dbscan_plot(clusters_df,color_clusters)

def dbscanplot2(dataset, epsilon, minsamples):
    dbscan_data = dataset[[0,1,2,3,4,5,6,7,8]]

    #return dbscan_data
    model1 = DBSCAN(eps=epsilon, min_samples=minsamples, metric='euclidean').\
    fit(dbscan_data)
    #return model1
    outliers_df = dataset[model1.labels_ == -1]
    clusters_df = dataset[model1.labels_ != -1]
    colors = model1.labels_
    color_clusters = colors[colors != -1]
    color_outliers = 'white'
    clusters1 = Counter(model1.labels_)
    return model1.labels_

```

```
[178]: yo =dbscanplot1(ZSHUT,.66,100)
```

```
Counter({0: 54416, 1: 2334, 2: 887, -1: 363})
number of clusters: 3
```

```
[167]: ssa =dbscanplot1(ZSHUT,.46,10)
```

```
Counter({0: 54398, 1: 2324, 2: 885, -1: 332, 4: 19, 3: 17, 5: 15, 6: 10})
number of clusters: 7
```

```
[164]: #(Counter(yo1[0])+Counter(yo[1])+Counter(yo[2])+Counter(yo[3])+Counter(yo[-1]))
Counter(yo1)
```

```
[164]: Counter({0: 54470, 1: 2353, -1: 252, 2: 887, 3: 38})
```

```
[180]: 252/5800
```

```
[180]: 0.043448275862068966
```

```
[174]: yo1 =dbscanplot2(ZSHUT,.66,10)
```

```
[175]: clusters11 = Counter(yo1)
clusters11
```

```
[175]: Counter({0: 54470, 1: 2353, -1: 252, 2: 887, 3: 38})
```

```
[176]: ze1 = ZSHUT
ze1[10]=yo1
ze22 = ze1[ze1[10] != -1]

truer = ze22[9]
mine = ze22[10]
purity_score(truer,mine)
#np.sum(1/55348)
```

```
[176]: 0.8432153494493316
```

```
[120]: contingency_matrix = metrics.cluster.contingency_matrix(truer, mine)
#return np.sum(np.amax(contingency_matrix, axis=0)) / np.
↳sum(contingency_matrix)
np.sum(np.amax(contingency_matrix, axis=0)/np.sum(contingency_matrix))
```

```
[120]: 0.9325662054264966
```

```
[ ]: ###start of search procedure###
yo =dbscanplot1(ZSHUT,.66,10) #gets the number of clusters and outlier
↳percentage
yo1 =dbscanplot2(ZSHUT,.66,10) #gets the cluster_df which holds labels for the
↳clusters

ze1 = ZSHUT #inititalize the ZSHUT dataset
ze1[10]=yo1 #add the clusters_df to our zshut dataset
ze22 = ze1[ze1[10] != -1] #wherever the cluster is labeled -1 aka an outlier we
↳will remove from the dataset

truer = ze22[9] #set original cluster labels as y_true
mine = ze22[10] #set the new labels we obtained as y_pred
purity_score(truer,mine) #run our purity score function to get the purity score
↳of our sclustering
```

```
[6]: shuttle = df2[[0,1,2,3,4,5,6,7,8]]
shuttle
```



```
[6]:
```

	0	1	2	3	4	5	6	7	8
0	50	21	77	0	28	0	27	48	22
1	55	0	92	0	0	26	36	92	56
2	53	0	82	0	52	-5	29	30	2
3	37	0	76	0	28	18	40	48	8
4	37	0	79	0	34	-26	43	46	2
...
57995	80	0	84	0	-36	-29	4	120	116
57996	55	0	81	0	-20	25	26	102	76
57997	55	0	77	0	12	-22	22	65	42
57998	37	0	103	0	18	-16	66	85	20
57999	56	2	98	0	52	1	42	46	4

[58000 rows x 9 columns]

```
[ ]: ###start of search procedure###
yo =dbscanplot1(shuttle,.66,10) #gets the number of clusters and outlier
    ↳percentage
yo1 =dbscanplot2(shuttle,.66,10) #gets the cluster_df which holds labels for
    ↳the clusters

ze1 = shuttle #initilize the ZSHUT dataset
ze1[10]=yo1 #add the clusters_df to our zshut dataset
ze22 = ze1[ze1[10] != -1] #wherever the cluster is labeled -1 aka an outlier we
    ↳will remove from the dataset

truer = ze22[9] #set original cluster labels as y_true
mine = ze22[10] #set the new labels we obtained as y_pred
purity_score(truer,mine) #run our purity score function to get the purity score
    ↳of our sclustering
```

```
[10]: yo =dbscanplot1(shuttle,10,10) #gets the number of clusters and outlier
    ↳percentage
```

Counter({0: 54172, 1: 2295, 2: 882, -1: 615, 3: 24, 4: 12})
number of clusters: 5

```
[11]: yo1 =dbscanplot2(shuttle,10,10) #gets the cluster_df which holds labels for the
    ↳clusters
```

```
[19]: ze1 = df2 #initilize the ZSHUT dataset
ze1[10]=yo1 #add the clusters_df to our zshut dataset
ze22 = ze1[ze1[10] != -1] #wherever the cluster is labeled -1 aka an outlier we
    ↳will remove from the dataset
truer = ze22[9] #set original cluster labels as y_true
mine = ze22[10] #set the new labels we obtained as y_pred
```

```
purity_score(truer,mine) #run our purity score function to get the purity score␣  
↪of our sclustering
```

```
[19]: 0.84522087653568
```